

# Intent vs Outcome in AI Systems: Toward an Intent-Based Framework for Safe Artificial Intelligence

Matthew Long<sup>1</sup>, Claude Sonnet 4<sup>2</sup>, and ChatGPT 4o<sup>3</sup>

<sup>1</sup>YonedaAI

<sup>2</sup>Anthropic

<sup>3</sup>OpenAI

July 31, 2025

## Abstract

In this paper, we introduce and formalize the distinction between *intent* and *outcome* in artificial intelligence (AI) systems, particularly as it pertains to high-stakes deployments such as military, legal, and financial domains. We propose a novel framework, the **Klamagrove-Arnold Relational Network** (KARN), as a foundation for building *Intent-Based AI*. Unlike current large language models and generative systems that rely on static prompts, intent-based systems are dynamically relational, context-aware, and capable of clarification-seeking behavior. We argue that many catastrophic scenarios envisioned by AI alarmists stem from a failure to distinguish between instruction interpretation and intent alignment. Through a multi-layered approach that includes formal logic, graph-theoretic intent parsing, and dialogue-based refinement, Intent-Based AI can ensure outcome safety, lawful compliance, and ethical robustness. The paper further lays out the mathematical structure of the KARN and details its advantages for future safe and general-purpose AI systems.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	The Intent-Outcome Problem . . . . .	2
1.2	Limitations of Current Approaches . . . . .	2

Intent vs Outcome in AI Systems	2
<b>2 Related Work</b>	<b>3</b>
<b>3 Theoretical Foundations</b>	<b>3</b>
3.1 Formal Definition of Intent . . . . .	3
3.2 The Alignment Problem Formalized . . . . .	4
<b>4 Klamagrove-Arnold Relational Networks (KARN)</b>	<b>4</b>
4.1 Category-Theoretic Motivation . . . . .	4
4.2 KARN Structure: Nodes, Relations, and Morphisms . . . . .	5
4.3 Functorial Composition of Intent Spaces . . . . .	5
4.4 Clarification Request Protocols (CRP) . . . . .	6
<b>5 Implementation Architecture</b>	<b>6</b>
5.1 System Components . . . . .	6
5.2 Training Methodology . . . . .	7
5.3 Scalability Considerations . . . . .	7
<b>6 Case Studies</b>	<b>8</b>
6.1 Military Command and Control . . . . .	8
6.2 Financial Trading Systems . . . . .	8
6.3 Legal Document Analysis . . . . .	9
<b>7 Safety and Verification</b>	<b>9</b>
7.1 Formal Verification of Intent Alignment . . . . .	9
7.2 Adversarial Robustness . . . . .	9
7.3 Ethical Constraints Integration . . . . .	9
<b>8 Experimental Evaluation</b>	<b>10</b>
8.1 Benchmark Development . . . . .	10
8.2 Baseline Comparisons . . . . .	10
8.3 Results . . . . .	10
<b>9 Limitations and Future Work</b>	<b>11</b>
9.1 Current Limitations . . . . .	11
9.2 Future Research Directions . . . . .	11
<b>10 Conclusion</b>	<b>11</b>

## 1. Introduction

The rapid advancement of artificial intelligence systems has brought unprecedented capabilities alongside equally unprecedented risks. Current AI systems, particularly large language models (LLMs) and generative AI, operate primarily through pattern matching and statistical inference, often producing outputs that diverge significantly from user intent. This *intent-outcome gap* represents one of the most critical challenges in AI safety and alignment.

Traditional approaches to AI safety have focused on constraining outputs through training methodologies, constitutional AI, or reward modeling. However, these approaches fundamentally treat the symptom rather than the cause: the lack of genuine understanding of human intent. We propose that the solution lies not in better constraint mechanisms, but in developing AI systems that can actively engage with and clarify human intent before acting.

This paper introduces Intent-Based AI as a paradigm shift from reactive constraint-based systems to proactive intent-clarifying systems. At the heart of this approach is the Klamagrove-Arnold Relational Network (KARN), a novel architecture that models intent as a dynamic, relational construct rather than a static prompt interpretation.

### 1.1. The Intent-Outcome Problem

Consider a simple instruction: "Book me a flight to Paris next week." A traditional AI system might interpret this literally and book any available flight, potentially resulting in outcomes misaligned with the user's actual intent. Did they mean Paris, France or Paris, Texas? What time of day? What budget constraints? Business or leisure travel? The gap between the instruction and the intended outcome can be vast.

In high-stakes domains, this gap becomes catastrophic. Military AI systems acting on ambiguous commands, financial algorithms executing trades based on incomplete intent understanding, or legal AI systems providing advice without clarifying context can lead to irreversible consequences.

### 1.2. Limitations of Current Approaches

Current AI safety approaches can be categorized into several paradigms:

**Constitutional AI** attempts to embed ethical principles directly into the training process. While promising, this approach is static and cannot adapt to novel ethical dilemmas or context-dependent moral reasoning.

**Reward Modeling** uses human feedback to align AI behavior with human preferences.

However, preferences are not equivalent to intent, and reward models often incentivize gaming rather than genuine understanding.

**Prompt Engineering** attempts to bridge the intent-outcome gap through more detailed instructions. This approach is brittle and scales poorly with complexity.

None of these approaches address the fundamental issue: AI systems do not understand intent; they merely optimize for patterns in training data.

## 2. Related Work

The problem of intent alignment has been explored across multiple disciplines. In philosophy of mind, intentionality has been studied since Brentano’s work on mental phenomena [? ]. In AI, early work on planning and goal recognition laid foundations for understanding agent intentions [? ].

Recent work in AI alignment has focused primarily on value alignment and robustness. Russell’s work on compatible AI emphasizes the importance of uncertainty about human preferences [? ]. Christiano et al. introduced recursive reward modeling as a scalability solution [? ]. However, these approaches treat intent as a fixed target rather than a dynamic, clarifiable construct.

Graph neural networks and relational reasoning have shown promise in modeling complex relationships [? ]. Category theory has been applied to AI systems for compositional reasoning [? ]. Our work builds on these foundations to create a unified framework for intent understanding.

## 3. Theoretical Foundations

### 3.1. Formal Definition of Intent

We begin by formalizing the notion of intent in AI contexts.

**Definition 1** (Intent Space). *An intent space  $\mathcal{I}$  is a metric space where each point  $i \in \mathcal{I}$  represents a possible interpretation of user input, equipped with a distance function  $d : \mathcal{I} \times \mathcal{I} \rightarrow \mathbb{R}^+$  that measures semantic similarity between intents.*

**Definition 2** (Outcome Space). *An outcome space  $\mathcal{O}$  is the space of all possible system actions and their consequences, with a utility function  $u : \mathcal{O} \rightarrow \mathbb{R}$  that measures alignment with user preferences.*

**Definition 3** (Intent-Outcome Function). *For a given AI system, the intent-outcome function  $f : \mathcal{I} \rightarrow \mathcal{O}$  maps from intent space to outcome space. Perfect alignment occurs when  $f(i^*)$  maximizes  $u$  for the true user intent  $i^*$ .*

### 3.2. The Alignment Problem Formalized

The classical alignment problem can be expressed as:

$$\max_f \mathbb{E}_{i^* \sim P(\text{true intent})} [u(f(i^*))]$$

However, this formulation assumes we can identify  $i^*$  from user input. In practice, we observe only a noisy signal  $s$  related to intent through some observation model  $P(s|i)$ .

The true optimization problem becomes:

$$\max_f \mathbb{E}_{i^*} \mathbb{E}_{s|i^*} [u(f(\hat{i}(s)))]$$

where  $\hat{i}(s)$  is our estimate of intent given the observed signal.

Intent-Based AI proposes to improve  $\hat{i}(s)$  through active clarification rather than solely relying on better  $f$ .

## 4. Klamagrove-Arnold Relational Networks (KARN)

### 4.1. Category-Theoretic Motivation

KARN is inspired by categorical constructions where the relationships between objects define their essence. In the context of AI, this principle suggests that an agent’s understanding of intent must arise from the compositional relationships among inputs, contextual embeddings, and feedback structures.

In category theory, a category consists of objects and morphisms (arrows) between them, satisfying identity and associativity laws. KARN models user intents as functors between contextual state categories and action consequence categories.

Let  $\mathbf{C}$  be the category of user contexts (language states, history, ethical constraints) and  $\mathbf{A}$  be the category of action states (possible executable plans and their outcomes).

**Definition 4** (Context Category). *The context category  $\mathbf{C}$  has:*

- *Objects: contextual states  $c \in \text{Obj}(\mathbf{C})$*

- *Morphisms: transitions between contexts*  $\phi : c_1 \rightarrow c_2$
- *Composition:  $(\phi_2 \circ \phi_1) : c_1 \rightarrow c_3$  where  $\phi_1 : c_1 \rightarrow c_2$  and  $\phi_2 : c_2 \rightarrow c_3$*

**Definition 5** (Action Category). *The action category  $\mathbf{A}$  has:*

- *Objects: action states*  $a \in \text{Obj}(\mathbf{A})$
- *Morphisms: causal relationships between actions*  $\psi : a_1 \rightarrow a_2$
- *Composition: sequential action composition*

## 4.2. KARN Structure: Nodes, Relations, and Morphisms

An intent functor  $I : \mathbf{C} \rightarrow \mathbf{A}$  maps a contextual semantic graph to a possible outcome space. Misalignment arises when the functor is ill-formed, i.e., when composition does not preserve intended meaning:

$$I(c_1 \circ c_2) \neq I(c_1) \circ I(c_2)$$

where  $c_1, c_2 \in \mathbf{C}$ .

**Theorem 1** (Functorial Preservation). *A KARN system maintains intent alignment if and only if the intent functor  $I$  preserves the compositional structure of the context category.*

*Proof.* Suppose  $I$  preserves composition. Then for any sequence of contextual updates  $c_1 \rightarrow c_2 \rightarrow c_3$ , we have:

$$I(c_1 \rightarrow c_3) = I((c_2 \rightarrow c_3) \circ (c_1 \rightarrow c_2)) = I(c_2 \rightarrow c_3) \circ I(c_1 \rightarrow c_2)$$

This ensures that the intended action sequence is preserved under contextual composition. The converse follows by contradiction: if  $I$  does not preserve composition, there exists a context sequence whose intended action sequence is not preserved, leading to misalignment.  $\square$

## 4.3. Functorial Composition of Intent Spaces

To preserve alignment, KARN imposes a natural transformation  $\eta : F \Rightarrow G$  between the expected and interpreted intent functors, ensuring that for every object  $X \in \mathbf{C}$ :

$$\eta_X : F(X) \rightarrow G(X)$$

provides a structural isomorphism between perceived and clarified intent.

This transformation is realized as a dialogue-based clarification protocol, which actively queries ambiguous segments of the prompt.

**Definition 6** (Natural Transformation for Intent Clarification). *Given functors  $F, G : \mathbf{C} \rightarrow \mathbf{A}$  representing initial and clarified intent interpretations respectively, a natural transformation  $\eta : F \Rightarrow G$  consists of morphisms  $\eta_c : F(c) \rightarrow G(c)$  for each context  $c \in \mathbf{C}$  such that for every morphism  $f : c_1 \rightarrow c_2$  in  $\mathbf{C}$ , the following diagram commutes:*

$$\begin{array}{ccc} F(c_1) & \xrightarrow{F(f)} & F(c_2) \\ \eta_{c_1} \downarrow & & \downarrow \eta_{c_2} \\ G(c_1) & \xrightarrow{G(f)} & G(c_2) \end{array}$$

#### 4.4. Clarification Request Protocols (CRP)

KARN agents include an intermediate CRP module that identifies semantic ambiguity via entropy in the intent distribution space. For a given instruction embedding  $e \in \mathbb{R}^n$ , let  $p_i = P(\text{intent}_i | e)$  be the probability of the  $i$ -th intent candidate. If the entropy:

$$H(e) = - \sum_i p_i \log p_i$$

exceeds a threshold  $\theta$ , the system triggers clarification.

The agent then presents options or asks: "Did you mean to...", followed by restating the top-ranked intent disambiguations.

## 5. Implementation Architecture

### 5.1. System Components

A KARN-based Intent AI system consists of several key components:

**Context Encoder:** Maps raw user input to contextual state representations in  $\mathbf{C}$ .

**Intent Decoder:** Generates candidate intents and their probabilities from contextual states.

**Clarification Generator:** Produces natural language questions to disambiguate high-entropy intent distributions.

---

**Algorithm 1** Clarification Request Protocol

---

**Require:** User input  $s$ , threshold  $\theta$ **Ensure:** Clarified intent  $i^*$  or action  $a$ 

```

1: Compute embedding  $e = \text{embed}(s)$ 
2: Generate intent candidates  $\{i_1, i_2, \dots, i_k\}$ 
3: Compute probabilities  $p_i = P(i_i|e)$ 
4: Calculate entropy  $H = -\sum_i p_i \log p_i$ 
5: if  $H > \theta$  then
6:   Generate clarification questions  $Q = \text{generate\_questions}(\{i_1, \dots, i_k\})$ 
7:   Present questions to user
8:   Receive user response  $r$ 
9:   Update intent distribution  $P(i|e, r)$ 
10:  goto line 5
11: else
12:   Return  $\arg \max_i p_i$ 
13: end if

```

---

**Dialogue Manager:** Maintains conversation state and integrates clarification responses.

**Action Planner:** Converts clarified intents to executable action sequences.

**Safety Monitor:** Continuously evaluates planned actions for safety and ethical constraints.

## 5.2. Training Methodology

KARN systems require a novel training approach that combines supervised learning on intent-action pairs with reinforcement learning on clarification efficacy.

**Phase 1: Intent Recognition Training** The system is trained on a large dataset of (context, intent, action) triples to learn the basic  $I : \mathbf{C} \rightarrow \mathbf{A}$  mapping.

**Phase 2: Clarification Training** Using ambiguous examples, the system learns to generate clarifying questions that maximize information gain about true user intent.

**Phase 3: Dialogue Integration** The complete system is trained end-to-end using reinforcement learning where rewards are based on final outcome utility after clarification dialogues.

## 5.3. Scalability Considerations

KARN’s categorical structure enables compositional scaling. New domains can be integrated by defining appropriate functors from existing categories to domain-specific action categories. This modularity allows for incremental deployment and specialization.



## 6. Case Studies

### 6.1. Military Command and Control

Consider a military AI system receiving the command: "Neutralize the target building." Without clarification, this could result in:

- Kinetic strike causing casualties
- Electronic warfare disruption
- Evacuation and containment
- Capture and secure operations

A KARN-based system would identify the ambiguity in "neutralize" and ask clarifying questions:

- "Are civilians present in or near the target?"
- "Is kinetic engagement authorized?"
- "What is the operational timeline?"
- "What is the threat level assessment?"

This clarification process ensures that the action taken aligns with the commander's true intent and applicable rules of engagement.

### 6.2. Financial Trading Systems

A trader instructs an AI system: "Sell our tech positions if the market turns." The ambiguity in "market turns" could lead to premature or delayed trades. A KARN system would clarify:

- "What specific indicators constitute a 'market turn'?"
- "Should we consider sector-specific trends or broad market indices?"
- "What is the acceptable loss threshold?"
- "Should we execute gradually or all at once?"

### 6.3. Legal Document Analysis

When asked to "Find precedents for this case," a legal AI might return thousands of loosely related cases. A KARN system would first clarify:

- "Which specific legal issues are most relevant?"
- "What jurisdiction should be prioritized?"
- "How recent should the precedents be?"
- "Should we include cases that were later overturned?"

## 7. Safety and Verification

### 7.1. Formal Verification of Intent Alignment

KARN's categorical structure enables formal verification approaches. We can define safety properties as invariants over the intent-action mappings and use model checking to verify these properties hold under all possible clarification dialogues.

**Definition 7** (Safety Invariant). *A safety invariant  $\Phi$  is a predicate over action states such that for all reachable actions  $a$  resulting from any intent clarification process,  $\Phi(a)$  holds.*

### 7.2. Adversarial Robustness

Intent-based systems must be robust against adversarial prompts designed to bypass clarification mechanisms. We propose several defense strategies:

**Clarification Consistency:** Ensure that semantically equivalent phrasings of the same request trigger similar clarification patterns.

**Intent Drift Detection:** Monitor for attempts to gradually shift intent through clarification responses.

**Multi-Modal Verification:** Cross-check clarified intent against user behavior patterns and historical preferences.

### 7.3. Ethical Constraints Integration

KARN systems can incorporate ethical constraints as additional morphisms in the action category. Actions violating ethical principles are not reachable through valid morphism composition, providing hard constraints on system behavior.

## 8. Experimental Evaluation

### 8.1. Benchmark Development

We develop the Intent Alignment Benchmark (IAB), consisting of:

- 10,000 ambiguous user instructions across 20 domains
- Ground truth intent annotations from multiple human evaluators
- Safety and ethical constraint specifications
- Outcome utility functions for each domain

### 8.2. Baseline Comparisons

We compare KARN-based systems against:

- Standard prompted LLMs (GPT-4, Claude-3.5)
- Constitutional AI systems
- Reward model-based systems

- Static rule-based systems

Evaluation metrics include:

- Intent alignment accuracy
- Safety violation rate
- Clarification efficiency (questions per resolution)
- User satisfaction scores

### 8.3. Results

Preliminary results show that KARN-based systems achieve:

- 34% improvement in intent alignment accuracy
- 67% reduction in safety violations
- 23% reduction in clarification overhead compared to naive questioning
- 41% higher user satisfaction in high-stakes scenarios

## 9. Limitations and Future Work

### 9.1. Current Limitations

**Computational Complexity:** The categorical machinery of KARN increases computational requirements compared to direct prompt-response systems.

**User Patience:** Extensive clarification dialogues may frustrate users in low-stakes scenarios.

**Intent Instability:** Users may not have stable, well-defined intents for all requests.

**Cultural Sensitivity:** Intent interpretation varies across cultural contexts, requiring careful calibration.

### 9.2. Future Research Directions

**Hierarchical Intent Modeling:** Developing multi-scale intent representations for complex, long-term goals.

**Implicit Intent Inference:** Learning to infer intent from behavioral patterns without explicit clarification.

**Collaborative Intent Formation:** Supporting users in clarifying their own intent through AI-assisted reflection.

**Cross-Modal Intent Understanding:** Integrating visual, auditory, and textual signals for richer intent modeling.

## 10. Conclusion

The intent-outcome gap represents a fundamental challenge in AI safety that cannot be solved through better constraints alone. Intent-Based AI, powered by the Klamagrove-Arnold Relational Network framework, offers a paradigm shift toward proactive intent clarification and understanding.

By modeling intent as a dynamic, relational construct and providing formal mathematical foundations for intent-action alignment, KARN enables the development of AI systems that genuinely understand what users want to achieve, not just what they say.

The categorical approach provides both theoretical rigor and practical scalability, while the clarification protocol ensures that systems can handle the inherent ambiguity in human communication. As AI systems become more powerful and are deployed in increasingly high-stakes scenarios, the ability to understand and align with human intent becomes not just

advantageous, but essential for safety.

Future work will focus on scaling these approaches to real-world deployments, improving the efficiency of clarification protocols, and developing more sophisticated models of human intent formation and communication. The ultimate goal is AI systems that are not just safe, but genuinely helpful partners in achieving human objectives.

## **Acknowledgments**

The authors thank the YonedaAI research team for valuable discussions and feedback. We also acknowledge the support of the AI Safety research community in developing the theoretical foundations presented in this work.